

Applying the end-end ML steps to a different dataset. ¶

We will apply what we've learnt to another dataset (airbnb dataset). We will predict airbnb price based on other features.

[25 pts] Visualizing Data

[5 pts] Load the data + statistics

- load the dataset
- display the first few rows of the data
- drop the following columns: name, host_id, host_name, last_review
- display a summary of the statistics of the loaded data
- plot histograms for 3 features of your choice

```
In [34]: %matplotlib inline
import matplotlib # plotting library
import matplotlib.pyplot as plt

import os
import sys
import tarfile
import urllib

DATASET_PATH = os.path.join("datasets", "airbnb")

import pandas as pd

def load_airbnb_data(airbnb_path):
    csv_path = os.path.join(airbnb_path, "AB_NYC_2019.csv")
    return pd.read_csv(csv_path)

airbnb = load_airbnb_data(DATASET_PATH) #load the dataset
airbnb.head() #display first few rows of the data

#drop the name, host_id, host_name and last_review columns
airbnb = airbnb.drop(["name", "host_id", "host_name", "last_review"], axis = 1
)


print(airbnb.describe()) #summary of statistics of the data


#plot histograms for any 3 features
airbnb["latitude"].hist()
plt.title("latitude")
plt.show()
airbnb["longitude"].hist()
plt.title("longitude")
plt.show()
airbnb["price"].hist()
plt.title("price")
plt.show()
```

	id	latitude	longitude	price	minimum_nights
\					
count	4.889500e+04	48895.000000	48895.000000	48895.000000	48895.000000
mean	1.901714e+07	40.728949	-73.952170	152.720687	7.029962
std	1.098311e+07	0.054530	0.046157	240.154170	20.510550
min	2.539000e+03	40.499790	-74.244420	0.000000	1.000000
25%	9.471945e+06	40.690100	-73.983070	69.000000	1.000000
50%	1.967728e+07	40.723070	-73.955680	106.000000	3.000000
75%	2.915218e+07	40.763115	-73.936275	175.000000	5.000000
max	3.648724e+07	40.913060	-73.712990	10000.000000	1250.000000

	number_of_reviews	reviews_per_month	calculated_host_listings_count
\			
count	48895.000000	38843.000000	48895.000000
mean	23.274466	1.373221	7.143982
std	44.550582	1.680442	32.952519
min	0.000000	0.010000	1.000000
25%	1.000000	0.190000	1.000000
50%	5.000000	0.720000	1.000000
75%	24.000000	2.020000	2.000000
max	629.000000	58.500000	327.000000

	availability_365
count	48895.000000
mean	112.781327
std	131.622289
min	0.000000
25%	0.000000
50%	45.000000
75%	227.000000
max	365.000000

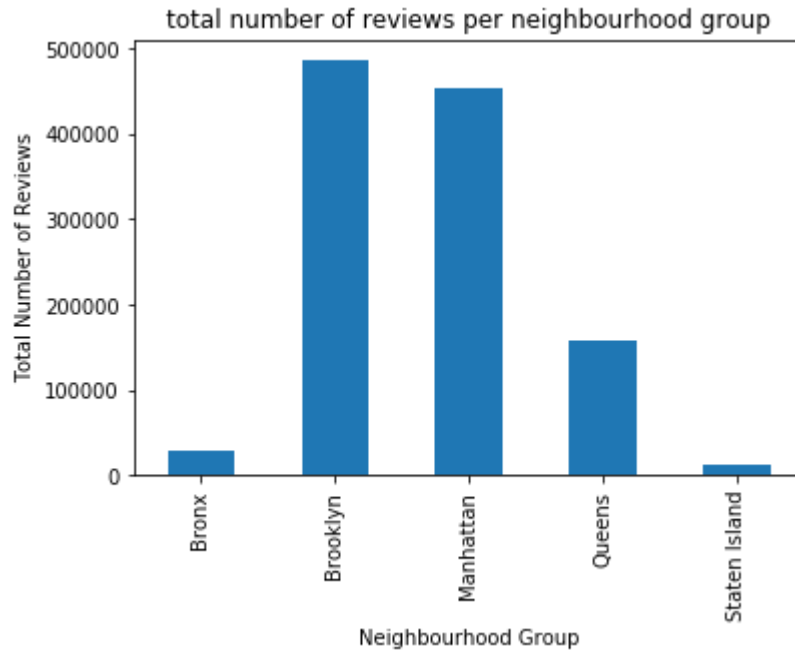




[5 pts] Plot total number_of_reviews per neighbourhood_group

```
In [35]: groups = airbnb.groupby("neighbourhood_group")["number_of_reviews"].sum().plot(
(kind="bar")
plt.ylabel("Total Number of Reviews")
plt.xlabel("Neighbourhood Group")
plt.title("total number of reviews per neighbourhood group")
```

Out[35]: Text(0.5, 1.0, 'total number of reviews per neighbourhood group')



[5 pts] Plot map of airbnbs throughout New York (if it gets too crowded take a subset of the data, and try to make it look nice if you can :)).

```

In [36]: images_path = os.path.join('.', "images/")
os.makedirs(images_path, exist_ok=True)
filename = "NYC.png"

#take a subset of the dataset so that the plot looks pretty
sample = airbnb.dropna() #drop rows that contain null values
#to take a sample of the subset I drop values from Brooklyn and Manhattan since they are overcrowded and make the map look ugly
airbnb_subset = sample.drop(sample[((sample["neighbourhood_group"] == "Brooklyn") |
                                     (sample["neighbourhood_group"] == "Manhattan")) &
                                     ((sample["number_of_reviews"] < 60) & (sample["availability_365"] < 180))].index)
print("Airbnb original set size: " + str(airbnb.size))
print("Airbnb_subset size: " + str(airbnb_subset.size))


import matplotlib.image as mpimg
nyc_img = mpimg.imread(os.path.join(images_path, filename),0)
ax = airbnb_subset.plot(kind="scatter", x="longitude", y="latitude", figsize=(10,7),alpha=0.1)

# overlay the NYC map on the plotted scatter plot
# note: plt.imshow still refers to the most recent figure
# that hasn't been plotted yet.
plt.imshow(nyc_img, extent=[-74.244, -73.713, 40.499, 40.913], alpha=0.5,
           cmap=plt.get_cmap("jet"))
plt.ylabel("Latitude", fontsize=14)
plt.xlabel("Longitude", fontsize=14)

save_fig("better_NYC_airbnb_plot")
plt.show()

```

Airbnb original set size: 586740
Airbnb_subset size: 213744
Saving figure better_NYC_airbnb_plot

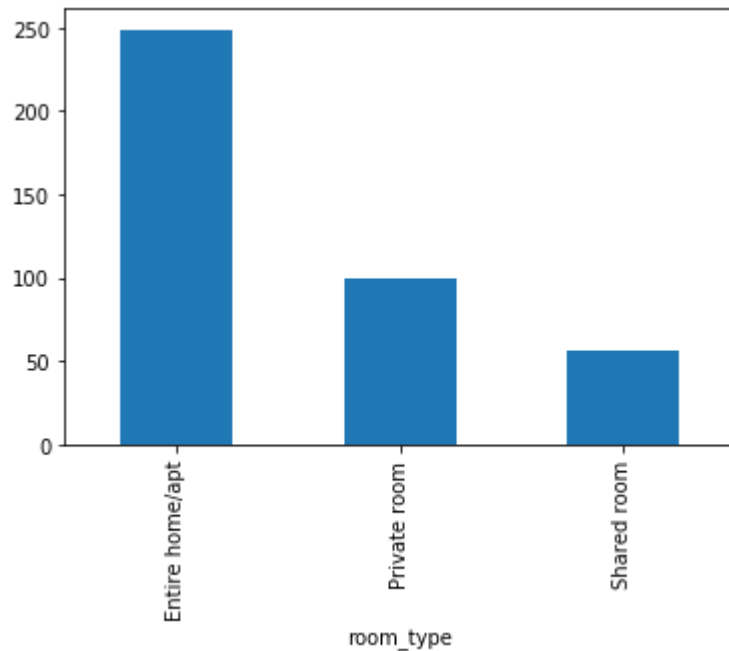


[5 pts] Plot average price of room types who have availability greater than 180 days.

```
In [37]: rooms_avail_180 = airbnb.drop(airbnb[(airbnb["availability_365"] <= 180)].index)

room_price = rooms_avail_180.groupby("room_type")["price"].sum()
room_count = rooms_avail_180.groupby("room_type")["price"].count()

av_room_price = (room_price/room_count).plot(kind="bar")
```



[5 pts] Plot correlation matrix


- which features have positive correlation?
- which features have negative correlation?


```
In [38]: #airbnb_corr_matrix = airbnb.corr()
#print(airbnb_corr_matrix)
airbnb_attributes = ["price", "minimum_nights", "number_of_reviews", "reviews_per_month", "availability_365"]
scatter_matrix(airbnb[airbnb_attributes], figsize=(18,12))
```

```

Out[38]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000141737701C8>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014173795D08>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000141737D2508>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000141741CB308>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000141744A3408
>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000141744DB508>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174514608>,
<matplotlib.axes._subplots.AxesSubplot object at 0x000001417454DC88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174557848>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174591A08
>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174677A48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000141749CFB08>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174A79C48>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174AAFD88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174AE8E88
>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000014174B21F88>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176C32048>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176C6A148>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176CA2248>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176CDA388
>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176D13488>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176D4A548>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176D85688>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176DBD788>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000014176DF4888
>]],
dtype=object)

```



[25 pts] Prepare the Data

[5 pts] Augment the dataframe with two other features which you think would be useful

```
In [39]: airbnb["minimum_cost"] = airbnb["price"]*airbnb["minimum_nights"]
airbnb["minimum_earnings_per_year"] = airbnb["reviews_per_month"]*airbnb["minimum_cost"]*12
```

[5 pts] Impute any missing feature with a method of your choice, and briefly discuss why you chose this imputation method

```
In [40]: from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy="median")
```

[10 pts] Code complete data pipeline using sklearn mixins

```
In [41]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.compose import ColumnTransformer

airbnb_num = airbnb.drop(columns=["neighbourhood", "neighbourhood_group", "room_type", "id"])

airbnb_num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler()),
])

airbnb_num_tr = airbnb_num_pipeline.fit_transform(airbnb_num)
airbnb_num_features = list(airbnb_num)
airbnb_cat_features = ["neighbourhood", "neighbourhood_group", "room_type"]

airbnb_full_pipeline = ColumnTransformer([
    ("num", airbnb_num_pipeline, airbnb_num_features),
    ("cat", OneHotEncoder(sparse=False, handle_unknown='ignore'), airbnb_cat_features),
])

airbnb_prepared = airbnb_full_pipeline.fit_transform(airbnb)
airbnb_prepared = pd.DataFrame(airbnb_prepared)
```

[5 pts] Set aside 20% of the data as test test (80% train, 20% test).

```
In [42]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(airbnb_prepared.drop(airbnb_prepared.columns[2], axis=1), airbnb_prepared[2], test_size=0.2, random_state=0)
```

[15 pts] Fit a model of your choice

The task is to predict the price, you could refer to the housing example on how to train and evaluate your model using MSE. Provide both test and train set MSE values.

```
In [43]: from sklearn.linear_model import LinearRegression

airbnb_lin_reg = LinearRegression()
airbnb_lin_reg.fit(X_train, y_train)

from sklearn.metrics import mean_squared_error

airbnb_preds = airbnb_lin_reg.predict(X_test)
mse = mean_squared_error(y_test, airbnb_preds)
print("Test mse value: ", mse)

airbnb_preds = airbnb_lin_reg.predict(X_train)
mse = mean_squared_error(y_train, airbnb_preds)
print("Train mse value: ", mse)

Test mse value:  0.6153495963161475
Train mse value:  0.6806895547681125
```